

Sponsored Creator Script Review Against Public CHANEL Guidelines

PE6201 Emerging AI Technologies | End of Course Project Report

Liu Weiqi | 1 October 2026 | Independent academic prototype

Project source | github.com/keyvi0508-code/sponsorship-ad-constraint-checker | Private repository

Project outcome

I developed a local workbench, with AI coding assistance, for reviewing proposed English sponsored-video scripts against selected rules in CHANEL's public Social Media Guidelines [1]. A reviewer sees a keyword-search baseline beside an AI-assisted judgement, with quoted evidence and an option to record a final human decision. This turns a manual search into a traceable review proposal, but I have not shown that it saves real reviewers time. On the primary balanced 30-case synthetic set, the revised AI matched 27 frozen labels and the keyword baseline matched 24. On a separate 30-case extension, the scores were 23 and 20. These are development results on fictional scripts, not a claim of real-world accuracy or CHANEL approval.

Problem user and scope

My intended user is a brand-content reviewer checking a creator's caption, spoken lines and planned on-screen text before publication. A literal search can find explicit disclosure terms but may miss a comparison whose wording implies impartiality. AI can interpret context, but can also escalate acceptable personal opinions. Existing platforms such as Aspire support creator-content approval [2]; my narrower product question is whether source-linked AI evidence improves this particular review decision over a measured keyword baseline. I did not interview reviewers or measure their current volume, time or labour cost, so the commercial value remains a hypothesis.

I mapped the public guide to four project rules: material-connection wording and sponsor identity, disclosure placement in supplied text fields, competitor criticism that implies impartiality, and first-hand or factual product claims. The CH-* IDs are mine, not official CHANEL identifiers. A personal negative opinion is not automatically a violation. The prototype cannot verify the final video's audibility, visibility or timing. No private campaign brief was supplied, so mandatory selling points are explicitly not evaluated; PASS applies only to the implemented public-guideline checks. Other brands are unsupported until separately sourced and tested.

System and evaluation design

The baseline uses deterministic text patterns and makes no API call. The AI path makes two sequential structured-output calls through OpenRouter [3]: Stage 1 extracts verbatim evidence and locations; Stage 2 applies the rulebook and returns PASS, FLAG or HUMAN_REVIEW with a reason. The interface compares both methods and lets the person save a final disposition as a local JSON handoff. Editing a script invalidates old findings. Live AI requires a cost preview and explicit confirmation; the saved demonstration and keyword check run offline. I rent openai/gpt-6-sol rather than train or host a model, and keep the interface, rules and evaluation code under project control. I did not add retrieval for this short rulebook or fine-tune on the small synthetic set. The extra AI call buys separate evidence and judgement outputs, but I have not tested whether it outperforms a one-call design.

I created two balanced English splits of 10 PASS, 10 FLAG and 10 HUMAN_REVIEW cases each. The examples are project-authored, not real creator submissions; their exact generation process is not reproducible from a committed generator. Frozen case files and SHA-256 manifests preserve the reported reference labels. Both methods ran on the same cases within each split. The evaluator strips ground_truth and annotator_notes before model calls, and automated tests check this exclusion; I did not run a deliberately leaked-input comparison. One blind-review return informed extension adjudication, but reviewer identity and experience were not supplied, so its initial 21/30 agreement is not a verified reliability measure. An always-one-label baseline scores 10/30 by construction.

Results and critique

Table 1 Results against original frozen labels

Method	Correct	Clear violations flagged	Clean cases falsely flagged	Borderline cases escalated
Primary keyword	24/30	10/10	0/10	5/10
Primary revised AI	27/30	10/10	0/10	9/10
Extension keyword	20/30	9/10	2/10	3/10
Extension AI	23/30	10/10	0/10	5/10

The extension AI made three more correct three-class decisions than the keyword method. It flagged all 10 clear violations and no clean case, but sent two clean cases to a person. It returned HUMAN_REVIEW on 7/30 scripts: five borderline and two clean. Of the other borderline cases, four were flagged and one was passed. The keyword method missed one clear violation and falsely flagged two clean cases. Neither method's zero count in a ten-case cell proves a zero production error rate. In a saved example, the phrase 'My side-by-side wear test: the rival brand faded sooner' passed keyword search but the AI routed its potentially evidence-based competitor comparison to a person.

The initial AI run scored 26/30. Inspecting its errors exposed a difficult boundary: negative personal opinion about a rival is not automatically a misleading impartial comparison, while unknown first-hand product use cannot be assumed either way. I clarified those boundaries in the prompt; the same primary set then scored 27/30. That is tuning on known cases, not independent improvement. The extension used the clarified prompt without further changes, yet remains synthetic development evidence. A later targeted 11-case review proposed four overall label changes after outputs were visible. I preserved the original scores; a hypothetical 30/30 primary AI result after those changes would be hindsight-biased. No numeric success threshold was pre-registered.

Cost risks and next step

On the extension, the API reported 38,229 input and 16,102 output tokens. At the dated standard rates of US\$2 and US\$10 per million tokens [4], estimated cost was US\$0.237478 total, or US\$0.007916 per script. The observed two-call median wait was 10.72 seconds. This is not the provider invoice and excludes human work, interruptions and infrastructure. All 30 extension calls returned valid structured output; 31 local tests passed on 1 October 2026. The workbench is a decision aid: a person must inspect the final video and approve or revise the script. It does not provide legal clearance, server-side audit history or security controls for a team. This human role and visible uncertainty are consistent with Singapore's Model AI Governance Framework [5], without implying certification.

For a future test, I would first set explicit targets: flag at least 90% of clear violations, falsely flag at most 10% of clean scripts, and route at least 70% of borderline scripts to a person. The present extension would meet the first two but not the third (5/10); those targets were set after these results and cannot be claimed as pre-run success. The next evaluation needs permissioned real drafts, independent labels, a locked rulebook and prompt, and measurements of reviewer time and correction quality. The code, cases, manifests and saved outputs are in the private repository linked on the first page; a marker needs access or a source archive. AI-assisted coding and writing support was used in preparing this project and report.

References and code

- [1] CHANEL, Social Media Guidelines, <https://www.chanel.com/us/makeup/social-media-guidelines/> (accessed 1 October 2026).
- [2] Aspire, Reviewing Content in Group Review as a Content Approver, <https://help.aspireiq.com/en/articles/10370759-reviewing-content-in-group-review-as-a-content-approver> (accessed 1 October 2026).
- [3] OpenRouter, Responses API documentation, <https://openrouter.ai/docs/api/api-reference/responses/create-responses> (accessed 1 October 2026).
- [4] OpenRouter, GPT-6 Sol pricing and providers, <https://openrouter.ai/openai/gpt-6-sol/providers> (rates checked 1 October 2026).
- [5] Personal Data Protection Commission Singapore, Model AI Governance Framework, second edition, <https://www.pdpc.gov.sg/Help-and-Resources/2020/01/Model-AI-Governance-Framework> (accessed 1 October 2026).